Cargill Seiveright

Software Engineer

Shinjuku, Tokyo, Japan | coolgill12@gmail.com | +81 80-3853-9559
github.com/gill876 | linkedin.com/in/cargill-s-a074b3125

Profile

Software engineer in Tokyo working on production Ruby on Rails, with a research background in NLP and LLM-based data processing (M.Eng., NAIST). At GA technologies I triage production errors for a real-estate investment service, RENOSY; I brought Claude Code and Snowflake into that work, cutting investigations from 1–2 days to under an hour, and made the process available to the whole team. In personal projects I benchmarked vision models to design a receipt-reading budget app that I run in production and use every week. Aiming to be an engineer who runs AI/LLM products reliably in production. Works in Japanese with partner teams (JLPT N2).

Experience

App Engineer *2024/11 – Present*
GA technologies Co., Ltd. *Roppongi Ichome, Tokyo, Japan*
RENOSY backend team (2026/05 – Present, team of 15)

- Triage production and staging errors from Rollbar alerts for the backend behind the customer-facing apps (iOS, Android, web) of RENOSY, GA technologies' real-estate investment service; the backend integrates with 14 internal and external services across 5 engineering teams. Isolate whether the cause is a domain service, infrastructure, or the application, and investigate through CloudWatch, CloudTrail, ECS logs, and a Snowflake mirror.
- Cut error investigations from 1–2 days to under an hour by being the first on the team to connect the Snowflake CLI to Claude Code; packaged the investigation steps as a Claude Code skill and published it on the internal marketplace so anyone on the team can use it.
- Identified three production data inconsistencies, one of which no one had known about, and worked with the domain teams to correct them.
- Spotted signs of memory fragmentation in memory monitoring and proposed adopting a tool to isolate the cause.
- Restructured escalation reports as "action needed → impact → cause", removing the follow-up questions that used to come back on each escalation.
- Wrote and corrected OpenAPI specs for 30+ of the backend's 100+ endpoints, checking each against the implementation, as part of a company-wide effort to end duplicated API docs between providers and consumers; coordinated with teams including iOS and Android.
- Set up a release check that verifies new features on both mobile and web; migrated RSpec fixtures to FactoryBot and replaced `create` with `build` to shorten CI runs.

Internal learning & work-management platform (2024/11 – 2026/04, sole developer)

- Built a full-stack Rails app that tracks employee learning progress and work hours in one place, owning basic and detailed design, DB modeling, API/frontend implementation, and tests.
- Designed many-to-many models with strict business-rule validations; implemented session-based authentication and authorization for nested resources.
- Wrote 300+ lines of RSpec/FactoryBot tests covering boundary values and edge cases; kept the code maintainable with helper extraction and RuboCop; localized the UI with a Japanese locale file; automated DB backup/restore with shell scripts to cut dev-environment setup effort.

Software Developer *2021/03 – 2022/08*
Artuvic Limited *Kingston, Jamaica (Remote)*

- Led frontend and backend enhancements for a nationwide courier platform (AngularJS, Java Play Framework), including model design for a new system and CSV-driven API integrations.
- Designed and implemented role-based access control and security features; built internal admin tools for operations and user management.
- Wrote JUnit unit tests for new system components and managed releases via Sourcetree.

Research & Publications

- **Master's thesis:** Can Large Language Models Outperform Rule-based Dataset Converters for Taxonomic Entities? (NAIST, 2024). Designed and evaluated an LLM pipeline that converts taxon NER datasets from Standoff to CoNLL format, handling the entity-boundary offsets, stray whitespace, and mixed character encodings that break rule-based converters; GPT-4o plus post-processing reached F1 ≥ 0.99 on all three datasets.
- **Co-author (2nd author):** Distilling Named Entity Recognition Models for Endangered Species from Large Language Models (arXiv, 2024). Generated 3,600 synthetic NER and relation-extraction sentences about endangered species with GPT-4 to distill into small BERT models; responsible for prompt design for data generation, fact-checking against the IUCN Red List, and writing part of the paper.

Projects

- **Kechi+** (2026/10 – Present, used weekly): a budget app that reads Japanese receipts from a photo (items, prices, tax category) and compares unit prices by store and month. Planned it and made the design decisions; built and operated it with Claude Code.

  - Benchmarked 11 vision models on GMI Cloud against my own receipts for accuracy, cost, and speed; chose a fast model for first reads, with a stronger model re-reading only when the totals don't add up.
  - Trunk-based development; GitHub Actions deploys to Cloud Run after lint and tests pass, then smoke-tests the live site; keyless auth via Workload Identity Federation. Monitoring with Sentry (errors, traces, per-release tracking), Grafana dashboards, SLOs, and a runbook.
  - Python (FastAPI), Next.js/TypeScript, PostgreSQL (Cloud SQL), Cloud Run, Firebase Authentication.
- **Flagship** (2026): a lab for trunk-based development, feature flags, and canary releases. Planned it and made the design decisions; built and operated it with Claude Code.
  - Every merge to main deploys a new Cloud Run revision at 0% traffic, released through a 0% → 10% → 100% canary.
  - Feature flags built on Rails configuration without a gem: segment targeting (dev, alpha, beta), sticky per-user percentage rollout, and a kill switch.
  - Because old and new revisions serve side by side during a canary, CI uses oasdiff to block breaking OpenAPI changes.
  - 12 bots browse as users, and a dashboard shows which revision and flag state served each one; monitored with Sentry, Grafana, and SLOs. Rails 8.1, Cloud Run, Firestore, GitHub Actions (Brakeman, bundler-audit, RuboCop).
- **5th Agentic AI Hackathon with Google Cloud** (2026): planned and built AsobiAgent solo, an AI agent that plans group outings in LINE groups. Owned the planning and design decisions and directed Claude Code and Antigravity to implement it. Chose not to let the AI book restaurants by phone, given the corporate contract it requires and the burden on restaurants; a person places the call with one tap instead. Vertex AI (Gemini), Cloud Run, Firestore, Secret Manager, LINE (LIFF), Ekispert API, Hot Pepper API.
- Competed in the ACM ICPC as an undergraduate.

Internships

Research Intern, NLP *2023/09 (3 weeks)*
Fujitsu Ltd. *Kanagawa, Japan (Remote)*

- Surveyed tokenization and embedding methods for low-resource NLP domains, focused on endangered-species corpora, and presented technical recommendations to Fujitsu's NLP Research team.

Software Intern *2020/06 – 2020/08*
The Palisadoes Foundation – Calico Challenge *Kingston, Jamaica (Remote)*

- Implemented encrypted communication between a central API and distributed data-collection agents for an open-source monitoring system (Pattoo) in Python, and built a React/Tailwind CSS admin portal for managing users, agents, and data pipelines.

Education

Nara Institute of Science and Technology (NAIST) 2022/10 – 2024/09
Master of Engineering, Natural Language Processing *Nara, Japan*
MEXT (Monbukagakusho) Scholarship recipient, Japanese Embassy recommendation

The University of the West Indies, Mona 2018/08 – 2021/09
Bachelor of Science, Computer Science (Minor: Japanese) *Kingston, Jamaica*

Skills

Languages: Ruby, Python, TypeScript, Java, JavaScript, SQL
Frameworks: Ruby on Rails, FastAPI, Next.js, React, AngularJS, Flask
Cloud & Data: AWS (ECS, CloudWatch, CloudTrail), Google Cloud (Cloud Run, Cloud SQL, Firestore), Docker, PostgreSQL, Snowflake
Observability & CI: Sentry, Grafana, Rollbar, GitHub Actions
Practices/Tools: RSpec, FactoryBot, TDD, OpenAPI, Claude Code, trunk-based development, feature flags, canary releases, Git
Spoken Languages: English (native), Japanese (JLPT N2, business working proficiency)